


CHAPTER 1

Solving Problems by Searching



This chapter includes the following points:

- ✧ Definition of intelligent agents.
- ✧ Description of one kind of intelligent agent called a problem-solving agent.
- ✧ Providing precise definitions of problems in AI and illustrating these definitions with examples.
- ✧ Description of several search algorithms that can be used to solve these problems.

1.1 Introduction to intelligent agents

1.1.1 What is an agent?

An **agent** is anything that can be viewed as perceiving its environment through sensors and acting on that environment through actuators (see Figure 1.1).

For example, a software agent receives keystrokes, file contents, and network packets as sensory inputs and acts on the environment by displaying on the screen, writing files, and sending network packets.

A robotic agent might have cameras for sensors and motors for actuators.




Figure 1. 1 Agent interacting with its environment

The agent's behavior is described by the **agent function** that maps any given percept sequence to an action. We evaluate the behavior of an agent using a performance measure; thus, the agent acts so as to maximize the expected value of the performance measure.

1.1.2 Intelligent agent

The job of AI is to design an agent program that implements the agent function.

There exists a variety of basic agent-program designs:

- 1- **Simple reflex agent:** This kind of agents selects actions on the basis of the current percept, ignoring the rest of the percept history (responds directly to percepts). It acts according to a set of condition-action rules. The structure of a simple reflex agent is shown in Figure 1.2.




Figure 1. 2 Simple reflex agent

- 2- **Goal-based agent:** Intelligent agents are supposed to maximize their performance measure. Achieving this is sometimes simplified if the agent can adopt a goal and aim at satisfying it. The structure of a goal-based agent is shown Figure 1.3.




Figure 1. 3 Goal-based agent

1.2 Problem-solving agent

One kind of goal-based agent is called a problem-solving agent. A problem-solving agent has three phases:

- ✧ Problem formulation;
- ✧ Searching solution;
- ✧ Executing actions in the solution;

Example 1: Pathfinding Problem

Imagine an agent on holiday in a country **X**. Suppose this agent is currently in city **A** and he has a flight that leaves tomorrow from city **M**. The road map of the country **X** is given in Figure 1.4.

State: In which city is the agent, i.e., the current location of the agent.

Goal: The goal of the agent is to be in city **M** as soon as possible, i.e., finding the shortest path from city **A** to city **M**, according to the road map.




Figure 1. 4 Example 1 graph (Country X road map)

1.2.1 Problem Formulation

One of the dominant approaches to AI problem solving is to formulate a problem/task as a search in a state space. A **problem** can be defined formally by the following four components:

- ✧ **Initial state:** the state that the agent starts in. For example, the initial state for the agent in the pathfinding problem might be described as $In(A)$.
- ✧ **Successor function:** the function that changes the state of the agent.

Given a particular state s , $successorsFn(s)$ returns a set of $(action, successor)$ ordered pairs, where each $action$ is one of the legal actions in state s and each $successor$ is a state that can be reached

from s by applying the action. An alternative formulation uses a set of operators that can be applied to a state to generate successors. For example, from the state $In(A)$, the successor function for the pathfinding problem would return $\{(GoTo(B), In(B)), (GoTo(D), In(D)), (GoTo(E), In(E))\}$, where $GoTo(B)$, $GoTo(D)$ and $GoTo(E)$ are the legal actions in state $In(A)$, and $In(B)$, $In(D)$, $In(E)$ are the successors.

State space is defined by the set of all states reachable from the initial state. Hence, the initial state and the successor function together implicitly define the state space of the problem.

The state space forms a directed graph in which the nodes are states and the arcs between nodes are actions as shown in Figure 1.5. A path in the state space is a sequence of states connected by a sequence of actions.




Figure 1. 5 Search tree for the pathfinding problem

- ✧ **Goal test:** which determines whether a given state is a goal state. Sometimes there is an explicit set of possible goal states, and the test simply checks whether the given state is one of them. For example, the agent's goal in the pathfinding problem is the singleton set $\{In(M)\}$.

Sometimes the goal is specified by a property rather than a set of one or more states. For example, in chess, the goal is to reach a state called "checkmate," where the opponent's king is under attack and can't escape.

- ✧ **Path cost:** It's the function that assigns a numeric cost to each path. The problem-solving agent chooses a cost function that reflects its own performance.

We can assume that the cost of a path can be described as the sum of the costs of the individual actions along the path. The **step cost** of taking action a to go from state $s1$ to state $s2$ is denoted by $c(s1, a, s2)$. The step cost for the pathfinding problem is the distance between the adjacent cities. For example, $c(In(A), GoTo(B), In(B)) = 75$.

A **solution** to a problem is a sequence of actions from the initial state to a goal state. Solution quality is measured by the path cost function, and an **optimal solution** has the lowest path cost among all solutions. For example, in pathfinding problem, the optimal solution is to find the shortest path from city **A** to city **M**.

1.2.2 Example problems

Example (a): Vacuum Cleaner Problem

Suppose we have a house with two rooms and one vacuum cleaner, and there is dirt in both of the rooms. In this problem, the vacuum cleaner is the problem-solving agent. The starting point of the vacuum cleaner is one of the dirty rooms, and its goal is to clean up the entire area.

Problem formulation:

- ✧ **States:** the state is determined by the location of the agent (the agent is in one of two rooms) and the condition of the rooms (dirty or clean). Thus, there are 2 (possible agent locations) * 2^2 (possible room conditions) = 8 possible states.
- ✧ **Initial state:** Both rooms are dirty, and the vacuum cleaner is in one of these rooms (see Figure 1.6).




Figure 1. 6 Possible initial states for the vacuum cleaner problem

- ✧ **Successor function:** This function generates the possible states that result from applying the three legal actions:

L: The vacuum cleaner moves to the left

R: The vacuum cleaner moves to the right

S: The vacuum cleaner sucks the dirt

The complete state space is shown in Figure 1.7




Figure 1. 7 State space for the vacuum cleaner problem

- ✧ **Goal test:** This test checks whether both rooms are clean (see Figure 1.8).




Figure 1. 8 Goal states for the vacuum cleaner problem

- ✧ **Path cost:** The cost of each step is 1, so the path cost corresponds to the number of steps in the path.

Example (b): Traveling Salesman Problem (TSP)

Given a set of cities and the distance between every pair of cities, the TSP problem is to find the shortest possible route (with the minimum cost) that visits every city exactly once and returns to the starting point (see Figure 1.9).




Figure 1.9 TSP graph

Problem formulation:

- ✧ **States:** A state is defined by a sequence of one or more cities visited by the agent. Each state must include not only the agent's current location but also the set of cities it has visited.
- ✧ **Initial state:** The initial state is the state *A*.
- ✧ **Successor function:** Each legal action corresponds to a trip to an adjacent city, with the precondition that each city must be visited exactly once. Therefore, the successor function generates the possible states that result from applying these actions. An example of a path in the state space is shown in Figure 1.10.
- ✧ **Goal test:** The goal test checks whether the agent is in the state *A* and has visited all 8 cities with the lowest cost.
- ✧ **Path cost:** The cost of each step is the distance between the adjacent cities, so the path cost is the sum of the distances.




Figure 1.10 Search tree for the TSP

Example (c): Blocks World Problem

We have a table with three blocks from A to C, arranged as shown in Figure 1.11 (a). Our goal is to rearrange them as shown Figure 1.11 (b). A robot arm can move only one block at a time, whether it's on the surface or on top of another block.




Figure 1. 11 Example of the Blocks world problem

Problem formulation:

- ✧ **States:** We will use a form of predicate logic to represent the state. Given below is the list of predicates as well as their meanings:

$On(A, B)$: Block A is on block B.

$OnTable(A)$: Block A is on the table.

$Clear(A)$: Nothing is on top of block A.

$Holding(A)$: The arm is holding block A.

$ArmEmpty$: The arm is holding nothing.

A state will be represented as the logical conjunction (AND) of the listed predicates.

- ✧ **Initial state:** $OnTable(A) \wedge On(B, A) \wedge OnTable(C) \wedge Clear(B) \wedge Clear(C) \wedge ArmEmpty$ (see Figure 1.12).




Figure 1. 12 Initial state for the blocks world problem

- ✧ **Successor function:** The robot arm can perform four actions:

$Stack(A, B)$: stacking block A on top of block B.

$Unstack(A, B)$: unstack block A from on top of block B.

$PickUp(A)$: Pick up block A from the table.

$PutDown(A)$: Place block A on the table.

Each of these actions has specific preconditions, which are represented using predicates. The effects of these actions are represented using two lists: DELETE and ADD. The DELETE list includes the predicates which will cease to be true once the action is performed, while the ADD list includes the predicates which will become true. Therefore, the successor function generates possible states (the conjunction of predicates) according to the information in Table 1.1. Figure 1.13 shows an example of a search tree for the blocks world problem.

Table 1. 1 Blocks world successor function

Action	Precondition	DELETE	ADD
$Stack(A, B)$	$Clear(B) \wedge Holding(A)$	$Clear(B) \wedge Holding(A)$	$On(A, B) \wedge ArmEmpty$
$Unstack(A, B)$	$On(A, B) \wedge ArmEmpty \wedge Clear(A)$	$On(A, B) \wedge ArmEmpty$	$Clear(B) \wedge Holding(A)$
$PickUp(A)$	$OnTable(A) \wedge ArmEmpty \wedge Clear(A)$	$OnTable(A) \wedge ArmEmpty$	$Holding(A)$
$PutDown(A)$	$Holding(A)$	$Holding(A)$	$OnTable(A) \wedge ArmEmpty$




Figure 1.13 Search tree for the blocks world problem

- ✧ **Goal test:** The goal test checks whether the goal state, represented as $OnTable(C) \wedge Clear(B) \wedge On(B, C) \wedge Clear(A) \wedge Holding(A)$, has been reached (see Figure 1.14).




Figure 1.14 Goal state for the blocks world problem

- ✧ **Path cost:** The cost of each step is 1, so the path cost is the number of steps in the path.

1.3 Searching for solutions

Having formulated a problem, we now need to solve it. This is done by a search through the state space. Search techniques use an explicit search tree to explore the state space.

The building block of a search tree is the node. There are many ways to represent nodes, but we will assume that a node is a data structure with five components (see Figure 1.15):

- ✧ **State:** the state in the state space to which the node corresponds;
- ✧ **Parent-node:** the node in the search tree that generated this node;
- ✧ **Action:** the action that was applied to the parent to generate the node;
- ✧ **Path-cost:** the cost, usually denoted by g , of the path from the initial node;
- ✧ **Depth:** the number of steps along the path from the initial state.




Figure 1.15 Structure of a node in a search tree

1.3.1 General search algorithm

- The root of the search tree corresponds to the initial state;
- The first step is to test whether this initial state is a goal state. If it is not, we must explore other states;
- We expand the current state by applying the successor function, which generates a new set of states;
- We continue the process of choosing, testing, and expanding until we either reach a goal state (i.e., find a solution) or exhaust all available states.
- The choice of which state to expand is determined by the **search strategy/search algorithm**.
- To assess the effectiveness of a search algorithm; we consider both the search cost and the path cost of the solution found.

The pseudocode for the general tree search algorithm is presented in Figure 1.16.

Input: The problem to be solved
Output: A solution or failure
Function TreeSearch (problem)
Begin
 Initialize the *Open* list using the initial state of the problem
 While the *Open* list is not empty do
 Choose a leaf node and remove it from the *Open* list
 If the chosen node contains a goal state, then return the corresponding solution
 Expand the chosen node
 Add the resulting nodes to the *Open* list only if they are not already in the *Open* list
 Return failure
End

Figure 1.16 Pseudocode for the general tree search algorithm

The pseudocode for the general graph search algorithm is shown in Figure 1.17.

Input: The problem to be solved
Output: A solution or failure
Function GraphSearch (problem)
Begin
 Initialize the *Open* list using the initial state of problem

```

Initialize the Closed list to be empty
While the Open list is not empty do
    Choose a leaf node and remove it from the Open list
    If the chosen node contains a goal state, then return the corresponding solution
    Add the chosen node to the Closed list
    Expand the chosen node
    Add the resulting nodes to the Open list only if they are not already in the Open or Closed lists
Return failure
End

```

Figure 1. 17 Pseudocode for the general graph search algorithm

1.3.2 Measuring problem-solving performance

Search strategies are evaluated based on the following dimensions:

- ✧ **Completeness:** Does the algorithm always find a solution if one exists?
- ✧ **Optimality:** Does it always find a solution with the lowest path cost?
- ✧ **Time complexity:** How many steps does it require to find a solution?
- ✧ **Space complexity:** How much memory is needed to perform the search?

Time and space complexity are expressed in terms of three quantities (refer to Figure 1.18):

***b*:** The branching factor, which is the maximum number of successors of any node;

***d*:** The depth of the shallowest goal node;

***m*:** The maximum depth of the state space (which may be $+\infty$);




Figure 1. 18 Time and space complexity parameters